

# Long-run future impact evaluation



Benjamin Todd, 28th May 2014

Executive Director

---

## Introduction

The purpose of this document is to evaluate the historical impact of 80,000 Hours from a long-run future perspective.

It is part of our May 2014 [annual review](#), and builds on our general purpose impact evaluation, most importantly:

- [Review of program performance and costs.](#)
- [Plan change analysis and cost-effectiveness.](#)

Note, we only have permission to share some of the information in this document with potential donors who are long-run future concerned, so please don't pass it on without asking us.

## Summary

*What's our strategy for improving the long-run future?*

We think building the capacity of the long-run future focused community is one of the most effective ways to improve the long-run future and mitigate existence risk.

We add talented, altruistic graduates to the long-run future focused community by introducing them to the reasons in favor of focusing on helping the long-run future through our programs (online content, coaching, and the 80,000 Hours community). Our programs also help those already in the community to have more impact with their careers.

*Do we have proof of concept that our programs can add altruistic, talented, graduates to the long-run future community?*

Yes. In total, we've collected 107 significant plan changes ([definition](#)). 18 of these (17%) have been substantially influenced by us towards wanting to use their careers to help the long-run future.

We think this provides proof of concept that our online content, coaching and community can lead to more highly able young people joining the community, and using their careers to help the long-run future.

*Have we helped people in the community have more impact?*

In addition, among the 107 people who made significant plan changes, we estimate nine already thought the long-run future was important but changed their career plans due to 80,000 Hours.

*Does the value of the eighteen people added to the community, and nine people helped within the community, justify our historical costs?*

Our total historical costs were £147,000 of donations and 13 years of labour by the team. If the team had pursued earning to give instead of working at 80,000 Hours, we estimate we could donated an additional £233,000 to high impact charities.

Historically, for every £1 we've received from the long-run future focused community, we've roughly been able to raise a further £1 from outside the community.

This group of 27 is highly able - including a Oxford maths lecturer, several startup founders and a Marshall Scholar - so we think they are likely to add more resources to the long-run future community than these costs.

To make a quantitative lower-bound estimate, we can look at additional charitable donations by the group. The four largest donors in the group expect to have donated £1,600,000 within the next three years, of which they attribute £540,000 to 80,000 Hours.

When you also take account of the value of donations beyond 2016 and the impact of the 18 people not pursuing earning to give (which includes the two founding staff members of the Global Priorities Project), we think it's likely the value of our contributions to the long-run future community significantly outweigh our costs.

*Does 80,000 Hours overall justify our costs in terms of resources added to the long-run future community?*

- We think the value of career changes of the 27 people caused by our programs already justifies our costs.
- In addition, we've built a body of research that has the potential to change many more careers in the future.
- We've also had other significant positive effects. For instance, we helped to incubate and develop the Global Priorities Project, and we've been one of the main groups building the effective altruism community.
- Throughout this time, our main focus has been on gaining value of information and investing for future growth, rather than immediate impact. As explained in our [strategic review](#), we think we've made strong progress improving our business model.

Overall, we think the evidence is sufficient to justify further developing our programs and continuing to expand our reach into our entire target market. See our plans for the future in our [summary business plan](#). Over the next year, we plan to invest in developing an online careers guide, which has the potential to (i) change careers (ii) bring more people to the site (iii) deepen our knowledge of how to maximise the impact of career choices. Our aim is to have this guide in place in time to benefit from the publicity campaign accompanying the launch of Will MacAskill's book, *Effective Altruism*, which has an estimated budget of over £300,000.

## Table of contents

In the rest of this document, you can find more information about our strategy for building the long-run future focused community and on the ways we've had an impact so far.

Introduction .....	1
Summary.....	1
Table of contents.....	3
How might 80,000 Hours be an effective long-run future intervention?.....	4
Has 80,000 Hours added more people to the long-run future community?.....	6
How do our programs encourage our users to care about the long-run future?.....	6
How many plan changes in total?.....	6
What proportion of people who made plan changes became more concerned with the long-run future?.....	7
To what extent have we helped to improve the careers of people already long-run future concerned?.....	7
How do we help improve long-run future careers?.....	7
How many people have we helped? .....	8
Does the value of the people added to and helped within the long-run future community justify our costs? .....	8
What were the costs? .....	8
Where did the resources come from?.....	8
How does the value of the people added to the long-run future community compare to costs? .....	8
Other resources influenced .....	11
Value of our research into career choices to help the long-run future.....	11
How have we supported organisations within the long-run future community? .....	12
Global Priorities Project (GPP) .....	12
Future of Humanity Institute (FHI).....	12
Machine Intelligence Research Institute (MIRI) .....	13
Centre for the Study of Existential Risk (CSER) .....	13
How have we built the effective altruism community? .....	13
Conclusion .....	14
Appendix.....	15
Complete list of plan changers added to the long-run future community .....	15
Some members of the long-run future community we helped have more impact .....	27

# How might 80,000 Hours be an effective long-run future intervention?

## *Our model*

We provide advice to talented young people on how to make the most difference with their careers through online content, one-on-one coaching and building a community focused on effective altruist career choice. By doing this, we enable our users to take more high impact careers.

This has the potential to improve the long-run future through building the capacity of the long-run future concerned community. We think building the capacity of this community is among the best ways to help the long-run future, because:

- **It's an opportunity for leverage:** \$1 invested in community building can yield more than \$1 invested in the next best long-run future interventions.
- **It keeps options open:** We're unsure what the best long-run future interventions will be in the future, so it's important to keep our options open. A community that's responsive to the evidence will be able to take the best interventions as they are discovered.

## *In what ways do we build the long-run future community?*

Our programs (online content, coaching, community) build the long-run future community, through:

1. Encouraging more altruistic, talented, young people to care about the impacts of their actions on the long-run future, and switch into paths that do more to improve the long-run future.
2. Helping people who are already concerned about the long-run future to have more impact with their careers.

Besides the effects of our programs, we also support the long-run future community in three other ways:

3. Do research into how best to aid the long-run future through career choices.
4. Support the growth of organisations within the long-run future community. We do this by helping organisations within the long-run future community find staff, gain funding, and receive publicity.
5. [Developing the effective altruism community](#) (through building CEA, supporting other individuals, networking, shaping the concept etc.) which in turn will go on to support the long-run future community.

## *Our strategy to date*

80,000 Hours is in its 'discovery phase'. Our key goal in this phase is to identify a scalable model for building the long-run future community.

This means we think the majority of the value we have created lies in the *value of information* gained and resulting *potential for future growth*. We think our total market is over 100,000 individuals. To see more on our progress on how we've developed during our discovery phase, see our [strategic review](#).

So far, our key aims in terms of growing the far future community have been:

- Get *proof of concept* that our programs can grow the long-run future community.
- Add enough people to the long-run future community to justify our costs to the long-run future focused community.

We have not focused on maximizing the extent to which we add people to the long-run future focused community. Rather, we prefer to engage people more broadly in effective altruism, and promote concern for the long-run future when we're larger.

#### *Outline of the rest of this document*

We will:

1. Overview the evidence that we've added more people to the long-run future community, and ask whether this constitutes proof of concept that our programs can build the community.
2. Consider whether the value of adding these people to the community justifies our historical costs.
3. Overview the extent to which we've helped people in the long-run future community with their careers.
4. Sketch the value of our research into how to choose careers with the aim of helping the long-run future.
5. Overview how we've helped to build the long-run future community through building the effective altruism community and supporting other organisations.

## Has 80,000 Hours added more people to the long-run future community?

### How do our programs encourage our users to care about the long-run future?

- We encourage people to maximise the impact of their careers, which includes encouraging an open-minded and impact-driven approach to selecting causes. We think this is likely to result in more people becoming interested in causes that benefit the long-run future.
- In conversations with coachees and on our blog, we suggest using [a framework](#) to assess causes which explicitly contains an evaluation of the potential long-run effects. We present the [arguments for taking long-run effects into account](#) and [tell people about the different ways to improve the long-run future](#).
- We've released a [ranked list of causes](#), which lists building the effective altruism movement, doing prioritisation research, and working to mitigate global catastrophic risks at the top.
- We promote careers at [MIRI](#), [GiveWell](#), [FHI](#), the Future of Life Institute and CSER.
- We expose our users to smart, honest, thoughtful people who think that the long-run future is overwhelmingly important.

### How many plan changes in total?

In our [review of program performance](#), we overview the performance of our user channel from reach to behaviour change, in the period from when we were founded in 2011 to the end of 2013.

In brief, we've had 140,000 unique visits to the site, several thousand of these have become 'engaged users' (i.e. are subscribed to our social media or mailing list), **107 have made a 'significant plan change'**, the majority of these are following through to change their behaviour, and many are staying engaged with the effective altruism community.

The *significant plan change* is a metric we developed to track our impact. It is defined as follows:

*An individual has made a significant plan change if they say they have changed their credence in pursuing a certain career path, cause or next step by 20% or more; they attribute this change to 80,000 Hours, and there's a plausible story about how engaging with 80,000 Hours caused this change.*

For instance, if someone says they anticipate going to medical school with probability 55% and law school with probability 45%, then they read an article on the 80,000 Hours blog and switch to 75% med school and 25% law school that would be the minimum possible to count as a significant plan change.

We count the plan change as plausibly due to us if both of the following are true:

- The person tells us that they changed their plans and they think it was due to us.
- We can point to the new information we gave them that changed their mind.

We have performed an in-depth analysis of what these plan changes consist of and how generally valuable they are in our [plan change analysis](#).

## What proportion of people who made plan changes became more concerned with the long-run future?

We went through the list of 107 people, and identified those for whom we think we had played a significant role in making them more concerned with the long-run future. We contacted each of them, and asked them about their views on the long-run future, and how they arrived at these views.

In total, we counted 18 '*plan changers added to the long-run future community*' (17% of people who made significant plan changes). We define this as:

- Has made a *significant plan change*.
- Thinks the long-run future is overwhelmingly important.
- Say they were significantly influenced by 80,000 Hours in arriving at this view.
- Can point to the information given by and interactions with 80,000 Hours that made them more concerned with the long-run future.
- Are now working on activities that are plausibly good ways to help the long-run future.

We expect the percentage to rise over time, since we've observed that the longer people have been involved the more likely they are to appreciate the arguments in favour of the importance of the long-run future.

We've listed all the individual responses in the appendix.

## To what extent have we helped to improve the careers of people already long-run future concerned?

### How do we help improve long-run future careers?

We have several thousand regular blog readers and have given formal one-on-one coaching to over 125 people. Our online content and coaching provides the following benefits:

- We provide advice on the best strategies for gaining leverage in order to help the causes you support, such as the idea of [earning to give](#).
- We provide information about promising career paths that is difficult to find elsewhere, such as our [in-depth analysis of the potential influence of becoming a politician](#).
- We help people make a rational comparison of their individual career options, through our career framework and [planning process](#).
- We introduce our coachees to other people in our community who can help them with their careers.

## How many people have we helped?

According to our [impact survey](#), about one quarter of the people who have made a significant plan change have donated to, volunteered for, or worked at an organisation that works on mitigating existential risk, suggesting at least 27 of the people who made significant plan changes are aiming to help the long-run future.

Subtracting the 18 people added to the long-run future community, this suggests there's roughly an additional nine people (10%) who are in the long-run future community and have changed their career due to 80,000 Hours.

We provide three examples in the appendix.

## Does the value of the people added to and helped within the long-run future community justify our costs?

### What were the costs?

At the end of 2013, our total historical financial costs were £147,000. In addition, about 13 years of labour had been spent on 80,000 Hours by the team. We think the opportunity cost of this labour is likely to be larger than our financial costs. [Estimates](#) of how much our team could have donated if we pursued earning to give instead suggest an opportunity cost of £233,000 in forgone donations.

Assigning the both the financial costs and the foregone donations (£380,000 in total) entirely to adding people to the long-run future community would result in a figure of £21,000 per person added to the community (according to the above definition), or £14,000 per person added or helped.

### Where did the resources come from?

Historically, we've been able to raise over 50% of our budget from outside the far future community because we simultaneously appeal to people interested in other causes. If you expect this to continue, there may be potential for additional donations to leverage other resources from outside the long-run future community.

## How does the value of the people added to the long-run future community compare to costs?

### *Ability of the group*

The group of 27 is extremely able. It includes a maths lecturer at the University of Oxford, a graduate of Entrepreneur First (a successful UK based seed accelerator), and a World Economic Forum Young Global Shaper. The majority are graduates of Oxford, Cambridge or Ivy League



Universities. They are also highly altruistically motivated, and many say that making a difference is the main aim of their careers.

You can see what each individual is doing in the [appendix](#).

### *Expected donations*

One way to get a lower bound on the value of these changes is to look at the expected additional donations that will result from them. Six (33%) of the people added to the community are pursuing earning to give, and there's likely to be roughly another three among the people helped.

We performed an [in-depth analysis](#) of how much extra money people pursuing earning to give give to high impact charities due to 80,000 Hours. We found that the average person pursuing earning to give who has made a significant plan change intends to donate an additional £12,000 to high impact charities within the next three years due to 80,000 Hours.

However, the majority of these donations come from the largest couple of donors, and the long-run future concerned group is tilted towards the most high earning. Among the 28 people, Andrew Farmer, Matt Wage, Matt Gibb and Sam Bankman-Fried intend to donate an additional £1,400,000 within the next three years, of which they attribute £480,000 (35%) to 80,000 Hours. Matt Wage has already donated about £200,000 to CFAR, 80,000 Hours, CSER and others, of which he attributes about 30% to 80,000 Hours. The group expects to be donating significantly more beyond 2016, because their earnings are expected to rise rapidly.

**In total, this adds up to donations of £540,000 within the next three years.** This is 50% more than total financial and opportunity costs of £380,000 as of the end of 2013. However, we also need to take account of the following:

1. Career opportunity cost
  - a. The relevant individuals may have had a significant impact in the careers they would have pursued otherwise.
  - b. If you think long-run future focused work is particularly important, the opportunity cost is not significant, however, because the majority would not have contributed to the far future community otherwise.
2. Biases in reporting
  - a. The relevant individuals support 80,000 Hours, so are likely to be biased towards telling us what we want to hear. They may also be overconfident in projecting their future donations.
  - b. On the other hand, there are biases working in the opposite direction, so as not wanting to admit being influenced by another group.
  - c. Overall, we suspect the figures are likely to be moderately overstated.
3. Discounting
  - a. The resources will arrive in the future, so have a lower value compared to costs incurred today.
  - b. Depending on your discount rate, this could reduce the value of the donations by about 10-50%.

4. The relevant individuals may not put all of their resources into the long-run future community.
  - a. Although they all agree the long-run future is overwhelmingly important, some think it's not optimal to work directly on long-run future focused interventions.

In addition, 60% of the donations come from Matt Wage, who was persuaded by our team before our official founding.

**However, we haven't yet taken account of the value of donations beyond 2016.** This is difficult to estimate, because it's hard to know what the group would have done otherwise further into the future, and the chance of them dropping out of earning to give becomes much more significant. On the other hand, the people in the group have only just started their careers, so their earnings will grow several-fold over the next ten years. All six people pursuing earning to give are working in hedge funds or tech startups, so there's a reasonable probability that one will earn over \$10m within the next ten years.<sup>1</sup> Overall, we think the expected value of the donations beyond 2016 are likely to be similar or greater than those within the next three years.

**Taking into account the value of donations beyond 2016, we think it's likely the value of the donations from the group of 27 will justify our costs.**

#### *Other activities beyond donations*

The remaining 66% of the group of 28 plan to use their careers to improve the long-run future through working at or founding long-run future concerned organisations, and using their positions to advocate for long-run future interventions. Many explicitly compared their options with earning to give, and decided they could do more good through other means. We expect, therefore, that these people will have a similar, or even greater, impact than those pursuing earning to give.

Specifically, the group includes:

- Owen Cotton-Barratt, an Oxford maths lecturer, who is now working full time as a researcher at CEA's Global Priorities Project, which aims to prioritise causes from a long-run future perspective, and has already done significant work with policy makers.
- Niel Bowerman, a WEF Young Global Shaper, who is now also working on the Global Priorities Project, and is setting up a new project to do effective altruism outreach.
- An Oxford PPE graduate, now working as an assistant to an MP, who is hoping to build a career in politics.
- A Marshall Scholar who is doing a Masters in Philosophy at Oxford, and wants to use their academic career to research population ethics and existential risk.

---

<sup>1</sup>Three are working at a proprietary trading firm, where about 10% of new entrants become partners within ten years, and partners earn one to ten million dollars per year.

If one valued the impact of the people not pursuing earning to give similarly to those who are, it would roughly triple our total impact. So, we think if you take these people into account, we're fairly confident that the value of the 27 people justifies our costs.

### *Years of labour moved*

An alternative way to think about our impact is in terms of 'years of labour added to the long-run future community.' In our [full plan change analysis](#), we estimated that a significant plan change means about 0.8 years of labour will be reallocated from a conventional career path to a high impact one within the next four years. This would imply that the 18 long-run future concerned people created means about 14 years of highly able labour will be added to the long-run future community in the next four years. In addition, there will be 8 years of labour within the community that has moved from lower impact to higher impact paths.

We've spent about 13 years of labour on 80,000 Hours, much of which would have gone to the long-run future community otherwise, and received donations of £147,000. So, on this basis, it seems that we could expect the value of the long-run future people created to be roughly equal our costs over the next four years.

However, this ignores the impact of the group of 27 beyond the next four years. They all intend to continue to contribute to the long-run future community beyond this point. Even in the face of significant drop-outs after the next four years, therefore, the value of the group of 28 seems likely to justify our costs.

### **Other resources influenced**

We've only considered the value of people added to the long-run future community according to our definition. This only includes people who (i) we know about (ii) have been substantially and knowably influenced, and (iii) have already changed their plans. Our total impact in terms of bringing new people into the long-run future community is likely to be significant greater.

Moreover, we haven't factored in the value of career changes for people not focused on the long-run future focused community. You can find out more about these in our [general impact evaluation](#).

## **Value of our research into career choices to help the long-run future**

A significant proportion of our time has been spent building a body of knowledge about how to maximise the social impact of career choices from a long-run future focused perspective. This research has the potential to improve many more career decisions in the future.

Our research outputs include:

- A framework for comparing career options in terms of their long-run impact.

- An exploration of the best ways to gain leverage with your career, including the idea of [earning to give](#) (backed up by an analysis of which careers offer the best earning potential) and an analysis of the [expected influence of politicians](#).
- A rational process for making and improving your career plan, based on our [career model](#), including methods to [mitigate biases](#).
- An exploration of various strategic considerations, such as [the importance of keeping your options open](#).
- An overview of [which causes](#) might be best to support from a long-run future perspective.

We think the plan changes undertaken by our audience, as detailed above, are good evidence that this research is valuable. In the future, we plan to carry out an external research quality evaluation.

## How have we supported organisations within the long-run future community?

### Global Priorities Project (GPP)

- The two main people running GPP, Niel Bowerman and Owen CB, became concerned about the long-run future in part due to 80,000 Hours and took jobs at CEA in part due to our influence.
- GPP is being incubated within CEA, and we built CEA's incubator along with Giving What We Can.
- 80,000 Hours commissioned [research on the importance of cause prioritisation](#), which was used to raise funding and attract talent to the project, and helped to develop the initial proposals for the project.

Overall, we think it's fair to say that GPP would not exist without 80,000 Hours.

The aim of GPP is to research global priorities from a far future perspective and to use this research to influence policy makers. To this end, they have already written a policy paper on 'unprecedented technological risks' that has been well received at the Cabinet Office, have discussed its implications with senior policy advisors to the Prime Minister, and have contributed to the Chief Science Advisor's annual report.

### Future of Humanity Institute (FHI)

- We introduced the FHI to Kristian Ronn and James Ting-Edwards, who worked at the FHI providing administrative support and grant writing during.
- GPP is a collaboration with the FHI.

Sean O'Heigeartaigh, Manager of the FHI, says:

*"Despite its relative youth as an organisation, 80,000 Hours has provided substantial support for FHI's research. For instance, in a recent recruitment round for a project manager, 3 of our 6 shortlisted candidates mentioned having been motivated by articles on the 80,000 Hours website describing how effective management, fundraising and public engagement can greatly increase the impact of a research team. As a result, we had a field of excellent quality for work that is usually considered less 'prestigious', but can be of equal or even sometimes greater value, than that of a researcher. I anticipate further support as members interested in supporting high-impact research reach positions of influence and earning power."*

## Machine Intelligence Research Institute (MIRI)

We're in close contact with MIRI. Luke Muelhauser, the Executive Director of MIRI, says:

*"I've sought out and gotten useful advice from 80,000 Hours several times, especially on media strategy, organizational development, and research metrics."*

## Centre for the Study of Existential Risk (CSER)

- 25% of the initial seed funding for the first staff member, Sean, came from Matt Wage, who partly attributes his decision to pursue earning to give to us.
- We promoted CSER among Cambridge students by helping to organise CSER's [first public lecture](#), which was attended by about 300 people.

## How have we built the effective altruism community?

The ways in which we've contributed to building the effective altruism community are listed in our general-purpose overview, [what impact has 80,000 Hours had?](#) In brief, these are the ways we've built the effective altruism community:

- We've increased the size of the community. The majority of the 107 people who made significant plan changes were not taking an effective altruist approach to leading their careers before they started engaging with us. (More information is in our [general plan change analysis](#)).
- We've helped to build CEA, by covering 50% of its budget and doing key organisational set up. It has already incubated The Life You Can Save and Animal Charity Evaluators, it is currently incubating the Global Priorities Project and a project to do effective altruism outreach.
- By introducing members and career coachees to each other, we've improved networking within the community. Through this network, people have landed jobs at top hedge funds and founded new projects together (such as [dotimpact](#)).
- We helped to expand the focus of the community beyond global poverty and towards the long-run future.

- We were the first group to start using the term publicly on our website and have had [significant press coverage](#).

## Conclusion

We think our programs are already creating enough value through causing career plan changes to justify our costs and provide solid proof of concept that they work.

In addition, we think we've created significant value through body of research we've built, what we've learned about our business model, as well as our other contributions to building the effective altruism and long-run future focused communities. Overall, we think this means 80,000 Hours has been a highly effective use of resources from a long-run future focused perspective.

Now we have proof of concept, we're excited to further develop and scale up our programs, with the aim of maximising their impact.

---

## Appendix

### Complete list of plan changers added to the long-run future community

We wrote the following studies as part of our impact evaluation process, sent them to the individuals concerned by email and encouraged them to point out inaccuracies. We redrafted the studies until the individual was satisfied, and had them confirm the exact wording used. Direct quotes are indicated.

Where we don't have a full study available, or their views about the long-run future were unclear, we emailed them the following the questions:

- In your opinion, when choosing between causes, how important are the effects on the long-run future?
- In your opinion, how important is reducing x-risk as a cause?
- How did you come to these beliefs, and what role did 80,000 Hours play?

One was made anonymous. Two requested that we not share their study.

#### 1. Owen CB

**Background** Post-doc in mathematics.

**What were they planning before?** Intended to pursue an academic career in pure mathematics. Giving 10% to effective charities as a member of Giving What We Can.

**How did their plans change?** Owen became convinced to consider impact in choosing research topics. This made him less inclined to continue research in pure mathematics, preferring areas with important but neglected questions.

**What are they doing now?** He's doing prioritisation research at the CEA-FHI Global Priorities Project.

**How and why did the change come about?** Owen met the 80,000 Hours community through his involvement in Giving What We Can in Oxford. He was introduced to the case for basing career choice on expected impact by people in this community. Through further discussions, he became convinced to use his research skills to work directly on the most pressing questions.

80,000 Hours also contributed to founding the Global Priorities project where Owen now works. 80,000 Hours contributed to the initial idea, and wrote up the arguments in favor of the project, which helped the project to secure funding, some of which came from members of 80,000 Hours.

**In their own words:** He said we had a 'significant impact' on his beliefs because we introduced him to the relevant arguments for the first time. He adds "One of the major influences of 80,000 Hours was the idea of framing a career as part of what you can achieve for good in the world rather than just thinking of achieving good via donations."

### **What are their views on the long-run future?**

"It is very likely that the most important effects of our actions today will be their long term consequences. We are still in the process of developing tools to properly weigh how different actions today will affect the long term (in expectation). Historically there has probably been a reasonable amount of alignment between actions with good short-run effects and ones we should expect to have good long-term effects. However it is likely that we are reaching a point in history where a significant gap is opening up between these, for example as we come to a point where society can have significant effects on the risk of human extinction."

### **How did they arrive at these views?**

"Reading and discussions with people at FHI and CEA."

## **2. Edward Green**

**Background** Studied computer science at Warwick. Graduate of Entrepreneur First.

**What were they planning before?** Split between working as a programmer, freelancing as a programmer, or starting a startup. Interested in but had vague ideas about making a difference through reducing existential risk or broad interventions to support economic growth.

**How did their plans change, and what are they doing now?** Still considering the same three career paths, but became convinced that reducing the risks from artificial intelligence is the most important cause. Decided to donate 2.5-5% of their income. Currently working as a programmer at a consultancy, having been recommended for the role by other members of 80,000 Hours. Intends to donate to the Machine Intelligence Research Institute.

**How and why did the change come about?** Became convinced of artificial intelligence risk being important while interning at 80,000 Hours and being introduced to others in the community.

### **In their own words:**

Interning at 80k and being involved in the community did make me think through in more depth which causes I thought were most important. I came to the conclusion that AI research is probably the most important cause right now.



## Testimonial:

"If you're young, smart and want to have a major positive impact on the world - you really should speak to the team at 80,000 Hours."

He plans to donate £5,000 to MIRI over the next three years.

### 3. Jacob Nebel

**Background** Princeton - Philosophy, Class of 2013

**What were they planning before?** Intending to either go to law school or start a career in academic philosophy.

**How did their plans change, and what are they doing now?** Decided to focus on researching issues around global prioritisation and the risk of human extinction within law and philosophy, as well as advocate for the importance of these issues. He's now studying a Masters in Philosophy at Oxford on a Marshall Scholarship.

**How and why did the change come about?** He attended a talk by us in Princeton, then spent two months interning with us in Oxford, during which he had several in-depth discussions with our staff. Learned about the importance of the long-run future and ways to have impact within academic philosophy. We also gave him significant advice in writing his Marshall application.

#### In their own words:

*80,000 Hours changed how I looked at my career prospects. At first, I thought 80,000 Hours's message was just about earning to give, but Will persuaded me to forgo a lucrative internship at a hedge fund to do research for 80,000 Hours. I learned from this experience that I could have a great impact through research, but only if I think very carefully about how to maximize that impact. And that's just what I hope to do through philosophy. [through 80,000 Hours] I met amazing people, and it is important to me that I know I have a network of smart, supportive folks in the years ahead.*

### 4. Richard Batty

**Background** Oxford University - Human Sciences, graduated 2012

**What were they planning before?** Uncertain intention to go into public health as a researcher or NGO worker.

**How did their plans change?** He decided to improve his programming skills, and pursue jobs in start-ups with the aim of creating valuable products and earning to give.

**What are they doing now?** He's working at a start-up consultancy in London which aims to improve board decision making.

**How and why did the change come about?** Richard came to the original 80,000 Hours talk, and was immediately convinced. He then helped to found 80,000 Hours, interned with us, and received coaching. During this time, he researched a wide variety of careers. As the result of this, he considered a much broader range of sectors and considered doing earning to give rather than direct work.

**Other benefits** Richard met a friend through 80,000 Hours, who later provided Richard with loans to learn to program and helped him find jobs in tech start-ups.

**When choosing between causes, how important are the effects on the long-run future?**

"In principle, the long-run future is the most important thing for me, but in practice I'm unsure whether we have enough evidence about possible interventions to choose an intervention."

**How important is reducing existential risk as a cause?**

"Like the above, in principle it is very important but I am unsure about working on it in practice. It also might be more important to work on making the future broadly better rather than trying to specifically prevent extinction."

**What role has 80,000 Hours played in you coming to these beliefs?**

"I probably wouldn't have these beliefs without 80k. I would have come across Less Wrong without 80k but I don't think it would have been as persuasive as being involved with the 80k community and hearing about non-AI related x-risks and broad interventions for the long-run future."

## 5. Max Siegel

**What are their views on the long-run future?** Believes that aiming to reduce the risk of extinction and cause prioritisation are among the most important causes.

**How did they arrive at these views?** They were presented with the arguments in favour during an internship at 80,000 Hours and careers coaching.

## 6. Andrew Farmer

**Background** Major in Math at MIT.

**What were they planning before?** Identified as a utilitarian. Choosing between a career in software engineering or quantitative trading. Had job offers within both.

**How did their plans change?** Became more convinced of the importance of donations compared to the direct impact of working in software. Within finance, became more inclined to take the option with the highest discounted earnings, rather than the best lifestyle. Believes they've become more likely to stick to their altruistic aims.

**What are they doing now?** Decided to take a job in quantitative trading and donate over half of his income to GiveWell recommended charities and meta-charities. Highly concerned with choosing the career in which they can have the largest impact, including long-run effects.

**How and why did the change come about?** Was persuaded by the importance of donations compared to direct impact through discussion with a member of 80,000 Hours. Received career coaching and was given estimates of direct impact, which confirmed this impression. Was also encouraged to consider the discovery value, career capital and degree of fit with both options, and was told about crucial considerations around existential risk.

**How much does he intend to donate over the next 3 years?** £300k to GiveWell recommendations, GiveWell and effective altruist organisations. He attributes ~15% of his total impact to 80,000 Hours.

## 7. David John

**Background:** University of Stanford - Masters in Financial Mathematics, Class of 2010; went to UBC as an undergraduate at age 14; over 7 years experience as a quant in finance

**Situation before:** Looking to find a job within finance that would allow him to do more good directly. Cared about improving finance as a way to increase the long-run prosperity of humanity.

**How 80,000 Hours intervened:** Introduced him to Animal Charity Evaluators, where he volunteered. Long-term follower of our blog. Received 4 coaching sessions and written research on his options, which introduced him to global prioritisation as a cause and other crucial considerations concerning how to help the long-term future.

**How his plans changed:** He wants work in impact investing to build up a group of people who want to do global prioritisation.

**What he's doing now:** Looking for jobs in impact investing.

**Other support provided:** Our support helped him to find a job at a start-up impact investing fund.

**In his words:**

*I had been thinking about how to make a difference for some time, and when I encountered 80,000 Hours I found all these smart people who had been thinking about it for even longer. I really liked 80,000 Hours' quantitative approach to thinking about impact. Perhaps the most important thing I've learnt from 80,000 Hours, and it was Niel who emphasized this to me, was that the difference in effectiveness between the best interventions and the median intervention can be multiple orders of magnitude...*

**What are their views on the long-run future?** Believes it's highly important to help the long-run future. Supports cause prioritisation as one of the best long-run future interventions.

**How did they arrive at these views?** Initially highly concerned to help the long-run future and ran a blog about improving the valuation of long-run effects within finance. Found out about existential risk for the first time on the 80,000 Hours blog, and was introduced to the full arguments and the arguments for cause prioritisation during a case study with 80,000 Hours.

## 8. Roman Duda

**Background** Roman did the BPhil in Philosophy at Oxford, and was choosing between a wide range of careers.

**What were they planning before?** He was planning to do a master's in Psychology, with the view to do go into research designing interventions to increase people's happiness, productivity and decision-making.

**How did their plans change?** He decided to work at 80,000 Hours instead. He intends to pursue his career on the basis of where he can make the most difference.

**What are they doing now?** Roman is working full time at 80,000 Hours.

**How and why did the change come about?** Roman came to care more about making a difference with his career while an intern at 80,000 Hours. The plan change came about as a result of a case study done on Roman's situation by 80,000 Hours whilst Roman was an intern at 80k.

**What are their views on the long-run future?**

"I think the long-run future likely dominates and reducing the chance of existential risks is likely the most important work we can be doing. I think that doing more research into cause selection and which careers do the most good are among the most important things to be doing right now, both for working to reduce existentials risks and for clarifying our picture of which causes to be devoting marginal resources into."

**How did they arrive at these views?**

"My views on the long-run future were formed as a result of reading Luke Muelhauser's and Nick Bostrom's writing, and attending an AGI conference organised by the FHI in Oxford. The fact that people are 80,000 Hours also take these issues seriously has likely strengthened by belief that this is something that's important."

## 9. Jess Whittlestone

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

Very important. I'm pretty convinced that effects on the long-run future dominate when it comes to choosing between causes/giving opportunities.

**In your opinion, how important is reducing existential risk as a cause?**

It seems intuitively to me like xrisk is the most important problem for us to be working on. I'm hesitant to claim that I believe it is the most important cause because I still have a lot of uncertainty around this - especially relating to difficulties of measurement and prediction - and it seems plausible to me that the best way to impact the long-run future *could* be to work on more certain near-term causes like economic growth, improving decision making/predictions, or even poverty alleviation. Again, a lot of this uncertainty is down to not feeling like I've thought about the issue in enough depth to really assert myself as falling on one side or the other.

**How did you come to these beliefs, and what role did 80,000 Hours play?**

I only came across the idea of xrisk or the argument that we should choose causes based on their effects on the long-run future through 80k, so it played a pretty big role! When I first came across the idea, it made sense to me intuitively and also seemed exciting and neglected, but I wasn't evaluating it as a cause quite critically, or in comparison to other causes like poverty. Since I've got more involved in the EA movement, I've been made more aware of both arguments for xrisk/future generations as an important cause, and against it. I'd say that 80k has mostly made me aware of arguments in favour of xrisk, and so contributed to my general impression that it is a very important cause. My uncertainty is mostly due to there being others in the EA movement who are less convinced, and not having spent enough time exploring both sides to feel confident saying where my views fall.

## **10. Niel Bowerman**

**Background** University of Oxford, PhD candidate in Climate Physics. Climate campaigner. Named a "Young Global Shaper" by the World Economic Forum.

**What were they planning before?** Pursue a career in climate policy or be a science advisor to NGOs or governments.

**How did their plans change?** Initially switched to considering earning to give in finance, but after several months decided to work within effective altruist organisations. Intends to pursue the career in which he can make the most difference.

**What are they doing now?** Working at CEA on outreach, prioritisation research and policy development.

**How and why did the change come about?** Attended a Giving What We Can event in Oxford, and was introduced to the arguments in favor of earning to give, which immediately convinced

him to rethink his approach. This led to him becoming heavily involved in the community in Oxford, and helping to set up 80,000 Hours.

**What are their views on the long-run future?** Believes our impact on the long-run future is the most important consideration, and believes that reducing the risk of extinction is one of the most important causes. Is currently working on these issues by promoting effective altruism and prioritisation research.

**How did they arrive at these views?** Through discussion with people in the Giving What We Can and 80,000 Hours community in Oxford, which introduced him to the arguments for the first time and persuaded him of their force.

## 11. Markus Anderljung

**What were they planning before?** Working to fight climate change or global poverty through campaigning or social entrepreneurship.

**How did their plans change, and what are they doing now?** Decided that promoting effective altruism is the most important cause. Significantly increased how likely he is to found an effective altruist organisation.

**In their own words:**

My career session motivated me to think far more deeply about my career choices and gave me the information I needed to do so.

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

"Effects on the long-run future of any choice are likely the most important seeing as that is (hopefully) where most people live. Therefore, if I can have some reasonable confidence that some causes differ considerably in their effects on the long-run future, that will be a big decider."

**In your opinion, how important is reducing existential risk as a cause?**

"Reducing existential risk seems like a very important task. I am currently torn between thinking that existential risk and global poverty are the most important causes. The reason I am unsure of existential risk is that there seems to be so many uncertainties surrounding it, while I can have quite a bit of confidence in the effects work to reduce global poverty will have."

**How did you come to these beliefs, and what role did 80,000 Hours play?**

"I would say that 80k is the organisation that has had me think about existential risk and effects on the long-run future in a serious way. I cannot say that 80k was what convinced

me of the ethical case, but the organisation has moved me towards thinking that there might be tangible ways to act on that ethical case.”

## 12. Marek Duda

**Background** Originally from a liberal-arts background, Marek worked his way into a job at a bulge bracket investment bank in Mexico City.

**What were they planning before?** Marek was planning to pursue a career in Finance/Business. He was at an investment bank, and planning to eventually make the jump to the buy side, or maybe doing an MBA and going into a different part of the finance/business world. His main focus was money, mainly for the purpose of being rich when older. He didn't really think about making a large positive difference, in part due to a lack of interest in it, but also due to an ignorance of that even being a possibility.

**How did their plans change?** The first change after his exposure to 80,000 Hours was to greatly increase his interest in making a large positive difference in the world, and he thought he would do so by donating to the most effective causes from his corporate career. Eventually, he became convinced that to have the largest positive impact with his career, working directly within organisations focused on promoting doing good in an evidence-based, cause neutral way, was probably be the best thing for him to do. Marek changed his career plan radically and quickly; having only landed a job in finance the previous year, he decided to leave it to go work at the Centre for Effective Altruism.

**What are they doing now?** Marek started working at CEA in March 2014.

**How and why did the change come about?** The change came about as a result of one on one coaching with 80,000 Hours, reading the 80,000 Hours website, and discussion with other people at CEA as well as from his other/unrelated networks.

**What are their views on the long-run future?** Marek thinks that our impact on the long-run future is likely to be the most important consideration, although remains relatively agnostic on the best policy for negating x-risk until further research and discussion. He thinks the idea of complex and difficult to predict flow through effects obfuscates an obvious path considerably, although FAI research and global coordination improvement seem to stand out as the front runners.

**How did they arrive at these views?** Through reading Nick Bostrom's work, as well as Less Wrong and other related literature. 80,000 Hours didn't convince Marek of his views on the long-run

future, but it opened him to the idea that he might be able to try and have a direct impact himself, whereas before he saw it as more of a problem which it was not his place to work on.

### **13. Greg Lewis**

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

Probably the dominant consideration, given the (hopeful!) size of the future, and plausible causal links between actions now carrying forward to alter long-run outcomes.

**In your opinion, how important is reducing existential risk as a cause?**

I'd probably (pace Nick Beckstead) fall on the 'targeted' side of the broad versus targeted debate: I agree that what is good now is convergent with what will be good for the future, but I don't think the convergence is strong enough for there not to be underfunded ways of making the future go particularly better. X-risk seems to take a very large proportion of this space, and so seems very important (what particular avenues one should take in trying to limit X-risk is not something I have any confidence pronouncing upon).

**How did you come to these beliefs, and what role did 80,000 Hours play?**

80k probably had an ancillary role insofar as it put me in touch with the sorts of people who were interested in this, and I was persuaded by osmosis. I'm unclear how to estimate the counterfactual, but my best guess is that 80,000 Hours accounted for 20% of my shift.

### **14. Matt van der Merwe**

**What were they planning before?** Highly uncertain, but most likely journalism, teaching or non-profit work.

**How did their plans change, and what are they doing now?** Pursuing earning to give in consultancy or finance. 80,000 Hours was a significant factor in this decision, though they were already aware of some of the general ideas through Giving What We Can.

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

Very important

**In your opinion, how important is reducing existential risk as a cause?**



I'm very uncertain about this. With regards to direct long-run future impact I'm unsure about whether x-risk or trajectory-shaping or other things are more important. This makes me think the most important cause is research to determine which of these or other causes are most important.

**How did you come to these beliefs, and what role did 80,000 Hours play?**

I first heard about long-run future stuff while on CEA internship, eventually did some reading on the ethics, and have since been slowly trying to find out more about the empirical side of things. Assuming the influence of being amongst CEA-types counts as influence of 80k, then it has had a vital role. I almost definitely would not have been exposed to the cause if it were not for 80k.

## **15. Matt Gibb**

**Background:** Oxford University - PhD in Computational Biology 2012

**What were they planning before?** Student. Planning to work in start-ups.

**How did their plans change, and what are they doing now?** Decided to donate 33% of his income, and the value of the equity of any start-ups he creates, to effective charities. Plans to lead his career on the basis of where he can make the most difference.

Now he's the co-founder of a start-up, Dropkick.kr, which matches investors with crowdfunding projects. All of the founders have made a legally binding agreement to donate 33% of the proceeds of a future sale of the company.

**How and why did the change come about?** Matt came to an 80,000 Hours talk, was exposed to the idea of earning to give, and was immediately convinced. He started volunteering for 80,000 Hours as one of the founding 6 members.

**How much have they donated over the last three years?** £2-3k to AMF and SCI.

**How much do they expect to donate over the next three years?** £100k in expectation, though likely to be zero.

**When judging the value of different causes, how important is it to consider their impact on the long-run future?**

"I guess the value of future lives is just the same as present lives, although my uninformed prior for how much I can effect these lives is low."

**How important is reducing the risk of extinction as a cause?**

"Perhaps the most important cause I know of, I am just unsure as to how much effect I can have on it."

**What role did 80,000 Hours play in your attitudes towards the above, if any.**

"It is unlikely I would have considered it unless I had met all the smart people at 80k."

## 16. Sam Bankman Fried

**Background** MIT - Physics, graduates in 2014

**What were they planning before?** Student. Considering politics, journalism and academia. Highly concerned with making a difference.

**How did their plans change?** Became more in favor of earning to give, at least in the short-run, and more concerned about causes that benefit the long-run future.

**What are they doing now?** Completed an internship at a proprietary trading firm and intends to take up a job offer in the summer. Intends to donate all earnings he doesn't need to live on and do well in his job. Final year at MIT. He plans to lead his career on the basis of where he can make the most difference.

**How and why did the change come about?** Sam was convinced through discussion with 80,000 Hours staff; introduction to the 80,000 Hours community; and lectures on the core concepts. Participating in the 80,000 Hours community contributed to finding his finance internship.

**How much have they donated over the last three years?** None

**How much do they plan to donate over the next 3 years?** ~\$100k, though ~\$1mn per year from 3 years time if he doesn't burn out. He attributes 20% of these donations to 80,000 Hours.

**What are their views on the long-run future?** Believes that our impact on the long-run future is the most important factor.

**How did they arrive at these views?** Discussion with Will MacAskill introduced him to the arguments.

There are also two case studies we don't have permission to share.

Some members of the long-run future community we helped have more impact

## 1. Matt Wage

**Situation before:** Student. Considering a wide range of careers, with the top contenders being doctor and maths professor. Member of Giving What We Can, pledging 10% of his income to cost-effective charities.

**How 80,000 Hours intervened:** Through involvement with us and other groups, he started to consider earning to give or working directly in 'effective altruist' non-profits. We gave him extensive feedback, advice and discussion of his plans, in person and in writing, weighing up the relevant arguments and providing relevant information, over the course of months. We encouraged him to pursue earning to give, due to his unusually large earning potential. Note that the decision was made before the official founding of 80,000 Hours, but involved our pre-founding team.

**How his plans changed:** He decided to pursue earning to give, in part due to 80,000 Hours and in part due to advice from others in CEA. He plans to lead his career on the basis of where he can make the most difference.

**What he's doing now:** He is working for a proprietary trading firm pursuing earning to give. He agrees that this was the right decision. He continues to receive research from one of our trustees on where to donate, and we've introduced him to several other members of 80,000 Hours.

He has already donated about 200,000 and expects to donate 1,000,000 over the next three years. He attributes 30% of his donations to 80,000 Hours' pre-founding team. He donates to 80,000 Hours and the Centre for Applied Rationality among other organisations. He has also helped to convince a recently retired Partner of his company to support 80,000 Hours financially.

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

"Probably the most important factor."

**In your opinion, how important is reducing existential risk as a cause?**

"Probably the most important cause."

**How did you come to these beliefs, and what role did 80,000 Hours play?**

"Talking with effective altruists. 80,000 hours didn't really play a role."

## 2. Kristian Ronn

**Background:** Lund University - Maths and Philosophy double degree

**What they were doing before:** Student with no concrete intentions.

**How 80,000 Hours intervened:** Kristian interned with CEA, and we connected him to a project

management job at FHI. At FHI he helped establishing CSER by building their website, newsletter, created an actionable funding plan and contributed to its implementation. He worked for FHI July 2013 - May 2014.

**What they are doing now:** He is starting his own company in Sweden, called Meta Mind, with the aim of conducting high impact projects. Right now Meta Mind is working on a product called [Normative](#).

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

"Very important."

**In your opinion, how important is reducing existential risk as a cause?**

"It dominates pretty much everything else (if done right!)."

**How did you come to these beliefs?**

"Through the writings of Nick Bostrom"

### 3. Robbie Shade

**Background:** Oxford University - Robotics DPhil, Class of 2012

**Situation before:** Student. Undecided between academia and industry.

**How 80,000 Hours intervened:** Robbie came to the original 80,000 Hours talk on earning to give, and was immediately convinced. He volunteered to help 80,000 Hours build our website.

**How his plans changed:** He decided to work in industry rather than academia and pledge to donate 25% of his income rather than 10% (which he decided due to Giving What We Can). Decided to donate to 80,000 Hours.

**What he's doing now:** Earning to give by working as a Software Engineer for Google.

**Donations:** He has donated roughly \$12,000 over the last 1.5 years of employment, to 80,000 Hours, SENS Foundation and Against Malaria Foundation. He intends to donate \$50,000 over the next three years. I think I would attribute 50% to Giving What We Can, and 30-40% to 80k. I would be donating ~10% of what I am donating if it wasn't for the ideas I was exposed to be GWWC/80k.

**In your opinion, when choosing between causes, how important are the effects on the long-run future?**

"Very important."

**In your opinion, how important is reducing existential risk as a cause?**

"Very important, and hugely underrated - if we don't concern ourselves with x-risk reduction and the species is wiped out, then anything else we've done will have been of little value."

**How did you come to these beliefs, and what role did 80,000 Hours play?**

"Simple utilitarian calculation, thinking about extrapolation of numbers of sentient beings that can be affected - either through sophisticated simulations or just humans colonising beyond earth - if we survive there's going to be a lot more of us than there are now, and suffering/happiness seems like something that should be summed over all conscious creatures. 80,000 Hours didn't directly lead to these beliefs, but conversations with people involved have strengthened these ideas."